

Generalizability of Neural Networks with a Fourier Feature Embedding


Nithin Raghavan
Richard Hu
Alberto Checccone

Background

- Why do we measure the generalizability of a neural network?
 - We want to measure the training accuracy and testing accuracy over time.
 - Generalizability can be viewed as a complexity measure of data that one can use to predict the test accuracy of the learned neural network.
 - We can thereby give a clear bound on the evolution of certain classes of neural networks so that we can see how they evolve over time.
 - This can also measure the richness of the class of functions that a neural network can learn.
 - 3D shape regression, 2D image regression, CT, MRI, etc.
 - Occupancy networks

Background

- Why measure Neural Network generalizability?
 - Relates training accuracy to test accuracy
 - Less complex = more generalizable



Background

- How do we measure the generalizability of a neural network?
 - Empirical Rademacher complexity directly gives an upper bound on generalization error
 - Rademacher complexity can give us an easily verifiable measure that can differentiate between true labels and random labels.
 - Using the neural tangent kernel, we can obtain closed-form bounds for the Rademacher complexity over training and testing set for neural networks.

Background

- How is generalizability measured?
 - Rademacher complexity upper bounds generalization error.
 - Allows for NN optimization.

What is a neural tangent kernel?

- The Neural Tangent Kernel (NTK) is a kernel function that describes the evolution of artificial neural networks during training using gradient descent.
- The NTK tools additionally lend themselves useful when using a 1-Lipschitz loss function.

$$\Theta(x, y; \theta) = \sum_{p=1}^P \partial_{\theta_p} f(x; \theta) \partial_{\theta_p} f(y; \theta) .$$

Neural Networks and the Neural Tangent Kernel

- 2-layer Neural Networks reduce to kernelized ridge regression

$$k_{\text{NTK}}(\mathbf{x}_i, \mathbf{x}_j) = \mathbb{E}_{\theta \sim \mathcal{N}} \left\langle \frac{\partial f(\mathbf{x}_i; \theta)}{\partial \theta}, \frac{\partial f(\mathbf{x}_j; \theta)}{\partial \theta} \right\rangle$$

What is a neural tangent kernel?


- A kernel function $k(x, y)$ is a positive-definite function that measures the similarity between vectors x and y .
 - Dot product kernels are of the form $k(x, y) = \Phi(x) \cdot \Phi(y)$, where Φ is a lifting into some different feature space.
- A Gram matrix H is an n by n Hermitian matrix such that $H_{i,j} = k(x_i, x_j)$, where k is a kernel
 - n is the number of sample points.
- Overview:
 - Generalize bound on NN
 - NN at infinite epochs \rightarrow kernel RR with no ridge (where kernel = NTK)
 - Nithin paper \Rightarrow can use four yaf to modify NTK
 - Note that generalizability \rightarrow small values of projections of labels onto NTK eigenvectors
 - So we can create bound 2 with good B

Kernels

- The kernel function $k(x, y) := \Phi(x) \cdot \Phi(y)$ for some Φ
- The (Positive Definite) Gram matrix is $H_{i,j} = k(x_i, x_j)$

Coordinate-Based MLP

- MLPs with a ReLU activation function taking in points from \mathbb{R}^d
- Input is generally very low dimensional
- Want to learn labels



Fourier Feature embedding

- Coordinate-based multi-layer perceptrons
 - Feed-forward neural networks that only take in the coordinates of an input, instead of the entire input itself
- A Fourier Feature embedding is an embedding $\gamma(\mathbf{v})$ of the input to a coordinate-based multi-layer perceptron of the following form:

$$\gamma(\mathbf{v}) = [a_1 \cos(2\pi \mathbf{b}_1^T \mathbf{v}), a_1 \sin(2\pi \mathbf{b}_1^T \mathbf{v}), \dots, a_m \cos(2\pi \mathbf{b}_m^T \mathbf{v}), a_m \sin(2\pi \mathbf{b}_m^T \mathbf{v})]^T$$

Fourier Feature embedding

- Low feature space = MLPs can't approximate well
- Idea: lift features onto surface of hypersphere
- A Fourier Feature embedding $\gamma(\mathbf{v})$ defined below:
 - $\mathbf{B} = [\mathbf{b}_1, \dots, \mathbf{b}_m]$ and m are hyperparameters
 - a_i 's are constant with $\|a\| = 1$

$$\gamma(\mathbf{v}) = [a_1 \cos(2\pi \mathbf{b}_1^T \mathbf{v}), a_1 \sin(2\pi \mathbf{b}_1^T \mathbf{v}), \dots, a_m \cos(2\pi \mathbf{b}_m^T \mathbf{v}), a_m \sin(2\pi \mathbf{b}_m^T \mathbf{v})]$$

Why use a Fourier Feature embedding?

- MLPs with ReLU have a “wide” kernel that leads to over-smoothing.
- Necessary to learn high frequency functions
- Fourier feature embedding allows modification of NTK kernel matrix



What is a Fourier Feature embedding?

- Why use a Fourier Feature embedding?
 - Recent advances in NTK theory have shown that MLPs with ReLU have a “wide” kernel that leads to “over-smoothing”.
 - Useful in high-dimensional tasks in order to prevent overfitting In low dimensions
 - Such an embedding is necessary to learn high dimensional features
 - We would like to measure the generalizability of coordinate-based neural networks using the Fourier Feature embedding with a fixed sample size, but varying the number of allowed frequencies.

Basic Rademacher complexity bound

- Observation: norm of a vector passed through Fourier Feature embedding is bounded by the number of frequencies:

$$\|\gamma(x_i)\|_2 = \sqrt{\sum_{i=1}^r \cos^2(2\pi b_i^\top x_i) + \sum_{i=1}^r \sin^2(2\pi b_i^\top x_i)} = \sqrt{r}$$

- First idea: create weak bound on generalizability using only number of frequencies

Basic Rademacher complexity bound

- To conduct this kind of analysis, we consider the class of neural networks whose weights are bounded by some constant (usually the case in practice)

$$\mathcal{H}' \doteq \{f_{\Theta} : \|w\|_2 \leq B'_2, \|u_j\|_2 \leq B_2 \ \forall j = [m]\}$$

- Where the neural network itself is a 2-layer ReLU network defined as:

$$f_{\Theta} = \frac{1}{\sqrt{m}} \sum_{i=1}^m w_i \phi(u_i^{\top} x)$$

Basic Rademacher complexity bound

- After some analysis, we concluded that

$$\text{Rad}_S(\mathcal{H}') \leq 2B_2B'_2\sqrt{\frac{r}{n}}$$

Advancements in NTK theory

- Arora, et al. (2019)
 - Relates generalizability to NTK eigendecomposition

$$\sqrt{\frac{2\mathbf{y}^\top (\mathbf{H}^\infty)^{-1} \mathbf{y}}{n}}$$

$$\mathbf{y}^\top (H^\infty)^{-1} \mathbf{y} = \mathbf{y}^\top \left(\sum_{i=1}^n \frac{1}{\lambda_i} v_i^\top v_i \right) \mathbf{y} = \sum_{i=1}^n \frac{(v_i^\top \mathbf{y})^2}{\lambda_i}$$


Why are NTKs useful here?

- Arora, et al. (2019) found that the generalizability during training of a neural network follows this (where v_i refers to the eigenvectors of the NTK):

$$\sum_{i=1}^n \frac{(v_i^\top y)^2}{\lambda_i}$$

Relating Fourier Feature embedding to the NTK

- NTK matrix modifiable using Fourier Feature embedding
- How do we quantify this?



Why are NTKs useful here?

- This implies that if we modify the NTK of the neural network, we can directly control the eigenvalues of the NTK, and therefore control the generalizability of the network as a whole.
- Two aspects: More frequencies means that the network is able to quickly learn high dimensional information, but also changing the b -values implies that different frequencies in the training set can be learned faster.
- How do we quantify this?

Better Rademacher complexity bound

- Arora, et al. 2019 showed NTK for a two-layer neural network is as follows:

$$h_{\text{NTK}}(z) = \frac{z(\pi - \cos^{-1} z)}{2\pi}$$

- Which bounds the Rademacher complexity:

$$\mathcal{R}_S \left(\mathcal{F}_{R,B}^{\mathbf{W}^{(0)}, \mathbf{a}} \right) \leq \frac{B}{\sqrt{2n}} \left(1 + \left(\frac{2 \log \frac{2}{\delta}}{m} \right)^{1/4} \right) + \frac{2\sqrt{2}R^2\sqrt{m}}{\sqrt{\pi\kappa}} + R\sqrt{2 \log \frac{2}{\delta}},$$

Better Rademacher complexity bound

- We obtain the following derivation:

$$h_{\text{NTK}}(\gamma(x_i)^\top \gamma(x_j)) = \sum_{j=1}^r a_j^2 \cos(2\pi b_j(x_i - x_j)) \left(\frac{\pi - \cos^{-1} \left(\sum_{j=1}^r a_j^2 \cos(2\pi b_j(x_i - x_j)) \right)}{2\pi} \right)$$

Better Rademacher complexity bound

- Neural Tangent Kernel theory was used here to obtain a testing generalization bound of (where H^∞ is the NTK Gram matrix)

$$\sqrt{\frac{2y^\top (H^\infty)^{-1}y}{n}}$$

- Fourier Feature embedding reduces equation on previous slide to:

$$\sqrt{\frac{2}{n} \sum_{i=1}^n \sum_{j=1}^n y_i y_j h_{\text{NTK}}(\gamma(x_i)^\top \gamma(x_j))^{-1}}$$

Next steps: Measuring Empirical Generalizability

- Convert inputs to be uniformly sampled over $[0, 1)^d$ forms spherical convolution over input space
 - i.e., 2-D occupancy network, image memorization
- Implies eigenvectors of Gram matrix form DFT matrix

$$W = \frac{1}{\sqrt{N}} \begin{bmatrix} 1 & 1 & 1 & 1 & \dots & 1 \\ 1 & \omega & \omega^2 & \omega^3 & \dots & \omega^{N-1} \\ 1 & \omega^2 & \omega^4 & \omega^6 & \dots & \omega^{2(N-1)} \\ 1 & \omega^3 & \omega^6 & \omega^9 & \dots & \omega^{3(N-1)} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{N-1} & \omega^{2(N-1)} & \omega^{3(N-1)} & \dots & \omega^{(N-1)(N-1)} \end{bmatrix}$$

Next steps: Rademacher bound optimization

- Use experiments to calculate which b values achieve best bound
- Quantify importance of b values through occupancy networks

Next steps: Rademacher bound of RKHS norm

- Another idea: NNs trained using gradient descent are in the span of the training data in a reproducing kernel Hilbert space (RKHS) induced by the NTK.

Next steps: Rademacher bound of RKHS norm

- Li, et al. (2020) finds the complexity of functions evaluated with an RKHS-norm converging to kernel ridge regression
 - Much harder than the other path
- We can then apply Rademacher analysis on the kernel ridge regression problem using the RKHS norm, and using FF embedding

$$\|\hat{f}_l - \hat{f}_{l-1}\|_{\mathcal{H}}$$

$$\hat{\mathfrak{R}}_{v^l}(\mathcal{L}(\hat{f}_{l-1}, D; f_l, B_l)) \leq \frac{DR}{\sqrt{M}}.$$